

Mikhail Kogan

Senior Platform Engineer | AI Infrastructure | DevOps

About me

Platform Engineer with 10+ years shipping consumer products at scale. Recently built [CyberMem](#) — production-grade IaC platform automating multi-platform deployments. Proven track record: scaled teams 1 → 13 engineers, cut CI/CD cycles 2 weeks → 2 days, maintained 99.5–100% uptime. Combining product thinking with infrastructure expertise.

Experience

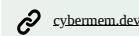
▼ CyberMem

Open Source Project: Production-grade Model Context Protocol (MCP) framework solving **AI context fragmentation across different LLM agents** (Claude, ChatGPT, Cursor, Perplexity). Features automated multi-platform deployment, **real-time observability**, and Zero-Trust networking via Tailscale. cybermem.dev / [GitHub](#) / [NPM](#)

Creator & Lead Developer: 2025 — Present

- Architected a **production-grade MCP server** orchestrating **shared memory** for different AI agents (Claude, ChatGPT, Cursor).
- Engineered a custom TypeScript CLI and IaC engine that **automates complex multi-platform deployment** (Mac, Raspberry Pi, Kubernetes) with a single command.
- Implemented remote access policy and observability**, enabling **scalable and zero-config production environments**.
- Automated, published and maintained** an open-source monorepo, achieving seamless multi-platform deployments and **100% CI/CD automation** via GitHub Actions.
- Monorepo Architecture: Built using TypeScript + CLI template engine, **achieving 100% CI/CD automation** via GitHub Actions.
- Zero-Trust Networking:** Implemented secure remote access and observability via Tailscale Funnel and Traefik SSL.
- Full observability stack: **Prometheus metrics + custom monitoring dashboard** for real-time tracking of AI agent operations.
- E2E Reliability: Architected a **robust testing framework using Playwright** to ensure mission-critical stability.

Details



cybermem.dev



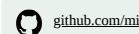
Tel Aviv, Israel



mikhailkogan17@gmail.com



linkedin.com/in/mikhail-kogan-platform



github.com/mikhailkogan17

▼ Centerya

A CRM platform powered by AI, designed for real estate agents in the US market.

📍 [Tel Aviv, Israel](#)

Founding Engineer & Mobile Infrastructure:

2024 — 2025

Centerya iOS App: <https://apple.co/3TuHEJB>. A mobile app for AI CRM with support of VoIP calls and Siri Intents.

- Infrastructure Ownership:** Implemented **real-time monitoring (Sentry)**, **centralized logging**, and **automated alerting** to achieve more than 99% crash-free sessions.
- Shipped weekly feature updates via **CI/CD Best Practices** using Xcode Cloud pipelines.
- Architected and solo-delivered AI-powered CRM app.
- Integrated Twilio SDK for VoIP calls with **AI-driven summaries & property suggestions**.
- Delivered iOS app with SwiftUI + SwiftPM → modular, scalable, optimized for **caching & startup performance**.

Personal Projects

Hearly — Real-time speech-to-text and translation app that helps users follow live conversations in multilingual environments, including group discussions and public settings. Built with SwiftUI, Azure Speech, and OpenAI Whisper; launched in 2025. Available in [hearly.app](#).

SoundPulseButton — open-source SwiftUI button component with haptic and audio feedback, featuring pulse and ripple animations. Listed on [Swift Package Index](#) / [GitHub](#); also used in Hearly.

Skills

Infrastructure & DevOps

- IaC:** Docker, Kubernetes, Helm, Ansible, Terraform basics, Traefik
- Programming Languages:** TypeScript, Python, Bash
- CI/CD:** GitLab CI, GitHub Actions, Xcode Cloud, Fastlane
- Observability:** Prometheus, Grafana, Sentry, Firebase, Vector, Amplitude
- Networking:** Tailscale, AWS, Zero-Trust Architecture

Mobile & Platform

- Swift, SwiftUI, Modern Concurrency, SDK distribution (XCFramework)
- Architecture: Modular, DDD, Feature Flags, TBD flow
- Testing: Unit, Snapshot, E2E (Playwright, Pytest)

Languages

Russian (native), English (fluent)

▼ Tinkoff Bank

The largest digital-only bank in Western Europe operated in the Russian market.

📍 [Russia](#)

Team Lead — Dolyame iOS platform: 2021 — 2023

Dolyame — Russia's first BNPL (Buy Now, Pay Later) service. Comparable to Klarna (Europe) and Afterpay (US).

- Launched the BNPL iOS app as a solo developer → scaled to **1M+ downloads and 300K MAU** in the first year.
- Automated **JIRA analytics pipeline** (Python + ChatGPT + JIRA API) to identify **team bottlenecks, workload distribution, and capacity planning**.
- Built & maintained an SDK (XCFramework) for partner integrations, supporting multiple versions and configs.
- Scaled the iOS app & SDK team from 1 → 13 engineers across multiple squads — enabled **faster delivery and product experiments**.
- Architected and configured UI Testing Workflow, allowing QA team to **reduce regression time by 70%** by easily automating their test cases.
- Scaled engineering team 1 → 13, establishing CI/CD standards and code review processes to enable rapid product growth.
- Transitioned team from Git Flow → TBD with feature flagging, **cutting release preparation time 4d → 1d** and allowing product team to **seamlessly toggle any feature on prod**.
- Maintained **99.5–100% crash-free rates** and high release stability across the codebase.

iOS Developer — Traffic Fines: 2018 — 2021

Traffic Fines App — a platform enabling users to conveniently pay their fines.

- Refactored critical payment & fine lookup modules in a legacy hybrid app, ensuring **99.7% crash-free sessions for 5M+ users**.
- Architected and maintained a shared library of **reusable UI components**, streamlining development across multiple banking apps.
- Established robust testing infrastructure (Unit, UI, Snapshot), achieving high coverage for financial transaction flows.

▼ GLOBUS Ltd.

📍 [Nizhniy Novgorod, Russia](#)

Junior iOS Developer: 2016 — 2017

- Developed several iOS apps, some achieving over 2K downloads.

Education

BSc in Computer Science & IT, Nizhniy Novgorod State University: 2013 — 2017